

Who's Playing Who? (من يلعب على مين؟)

A Lecture Script on AI — for students aged 15–18

Total runtime: ~45–50 minutes + Q&A

Stage directions are in *[brackets and italics]*. Lines you can say almost word-for-word are in plain text. "ASK" boxes are moments to stop and get the boys talking — don't skip them, they're what makes this land.

OPENING — The Hook (3–4 min)

[Walk in, no slide up yet, or just the title slide with the Arabic line on it.]

"Before I say anything about AI — I want you to play a game with me for ten seconds. Pull out your phone, open Instagram or TikTok, and look at the very first thing in your feed. Don't show me. Just look."

[Pause. Let them actually do it.]

"Now — raise your hand if that first video or post was *exactly* what you wanted to see right now."

ASK: *Get 2–3 hands. Ask one of them: "How did it know?"* Let them answer — they'll usually say something like "the algorithm" or "it just knows me."

"That's the whole lecture, actually. Something on the other side of that screen studied you — what you watch, how long you watch it, what you skip — and decided what to put in front of your eyes next. So here's the real question I want you walking out of this room asking yourself every single day:

Are you using AI — or is AI using you?

That's why this is called *Who's Playing Who*. Today we're going to open up the black box: what AI actually is, how it thinks (or doesn't), how it's already shaping your decisions without you noticing, how it's remixing music and movies, and — most importantly — how to use it so that *you* stay the one in control."

PART 1 — What Is AI, and Does It Actually Think? (10 min)

1.1 AI is not one thing

"When people say 'AI,' they're usually picturing a chatbot. But AI is already inside almost everything you touch. The Instagram feed we just looked at? AI. Netflix or Shahid telling you 'because you watched...'? AI. Face filters on Snapchat? AI. Google Maps predicting your traffic? AI."

ASK: "Give me one more example from your own phone — something you use every day that's secretly AI." *[Take 2–3 answers — FaceID, Spotify recommendations, autocorrect, etc. Validate each one.]*

"So AI isn't one app. It's a category — systems that take in data about you and the world, and use it to predict or generate something. Today we're going to focus on one specific type, because it's the one changing the world fastest right now: **Large Language Models**, or **LLMs**. This is the tech behind ChatGPT, Gemini, and Claude."

1.2 The three big names

[Slide: logos of OpenAI, Google, Anthropic]

"Three companies are racing each other on this technology:

- **OpenAI** — they built **ChatGPT**.
- **Google** (through its AI lab **DeepMind**) — they built **Gemini**.
- **Anthropic** — they built **Claude**.

These three are basically the Real Madrid, Man City, and Barcelona of the AI world right now — different teams, same competition, fighting for who builds the smartest model. Keep these three names in your head, because everything in the news about AI traces back to one of them."

1.3 How does an LLM actually "think"?

"Here's the part that surprises people. Open your WhatsApp keyboard right now and start typing 'I am going to...' and look at the three suggested words above your keyboard."

ASK: *Have a student read out what their keyboard suggests.*

"That little prediction bar is doing — in a tiny, dumb way — exactly what ChatGPT does, just at a massive scale. An LLM has read an almost unimaginable amount of text from the internet — books, articles, websites, conversations — and it learned the *statistical pattern* of which word

tends to follow which word. When you ask it a question, it's not 'understanding' your question the way your brain understands it. It is predicting, one word at a time, what the most likely next word should be, based on everything it has memorized. It's autocomplete on steroids."

1.4 So... does AI think?

[Show the screenshot: ChatGPT displaying "Thinking..." before answering.]

"You've probably seen this. You ask ChatGPT something hard, and before it answers, it shows you a little 'Thinking...' animation, sometimes even a paragraph of 'reasoning' before the final answer. Looks exactly like a person pausing to think, right?

Here's the truth: that's a performance. What's actually happening is the model is generating extra text *in the style of* reasoning — step-by-step working — because models that write out their steps tend to land on better answers, the same way you get fewer mistakes on a math exam when you show your work instead of just writing the final number. But it is not sitting there *having a thought* the way you are right now, listening to me and forming an opinion in your head. It has no awareness it exists. It has no goals of its own. It is mimicking the *shape* of thinking because that shape produces good outputs — it is not thinking."

ASK: "Why do you think a company would want their AI to *look* like it's thinking, even just for show?" *[Let 1–2 students guess — guide them toward: it builds trust, it feels more human, it makes you believe the answer more.]*

"Exactly. And that's our first hint of today's theme — something can be *designed* to feel trustworthy, without actually earning that trust. Keep that in your pocket. We'll come back to it."

PART 2 — Your Greatest Weapon: Critical Thinking (12 min)

"If AI doesn't really think — what does it mean that humans do? I'd argue the single greatest advantage you have over any AI model is **critical thinking**: the ability to question, doubt, and evaluate instead of just accepting what you're told."

2.1 Quick warm-up: spot what's wrong

"I'm going to read you some sentences. Your only job: tell me what's wrong with them. Don't overthink it."

[Read one at a time, pause for hands each time.]

1. *"This influencer takes a supplement and got abs in two weeks, so the supplement gave him abs."* — **ASK:** *"What's the flaw?" (He could've also been training and dieting — correlation isn't causation.)*
2. *"Everyone in my class likes this song, so it must be the best song ever made."* — **ASK:** *(Popularity isn't the same as quality — and "everyone" is a tiny, biased sample.)*
3. *"My phone's battery died right after I talked about buying new shoes, so my phone is definitely listening to me."* — **ASK:** *(Coincidence vs. evidence — no real proof of cause and effect.)*

"Good. That instinct you just used — pausing, asking 'wait, does that actually follow?' — is critical thinking. It's a muscle. And like any muscle, if you stop using it, it gets weak."

2.2 The "white paper" test — measuring creativity

"There's a famous test psychologists use to measure creativity. It's incredibly simple: they hand someone a blank sheet of paper and ask, **'How many different uses can you think of for this?'** Not 'what is paper used for' — how many *creative, unusual* uses can you invent in two minutes.

Let's try it right now. Two minutes, shout out as many uses for a sheet of paper as you can — the weirder the better."

ASK: *Run the 60–90 second brainstorm. Write a few on the board: paper airplane, fan, wrapping paper, fortune teller, makeshift cup, sandpaper substitute, origami, fly swatter...*

"Now here's the interesting part. Researchers actually ran this exact test on GPT-3, an early version of this same AI technology, and compared its answers to real humans. The AI could produce *a lot* of answers, fast — but on average, humans still beat it on originality and genuine surprise. AI is excellent at *remixing* what it has already seen. True creativity — the unexpected leap, the idea nobody has connected before — is still mostly a human thing. But here's the catch: that gap only stays in our favor if we keep practicing it. If you let AI generate every idea for you, every essay outline, every birthday card message, every project concept — that creative muscle, like the critical thinking one, starts to shrink."

2.3 When AI quietly damages your thinking

"This isn't just a theory. In 2025, researchers at MIT ran a real experiment. They split people into three groups writing essays: one group used ChatGPT, one used a normal search engine, and one used just their own brain — no tools at all. They measured brain activity while people wrote, and then tested how well people remembered or could explain their *own* essays afterward.

The ChatGPT group's brain activity was the weakest of the three. And when researchers later asked them to quote a sentence from the essay *they themselves had just 'written'* — most of them couldn't. They called this effect '**cognitive debt**': it feels free in the moment, exactly like ChatGPT did the thinking for you for free — but you're quietly borrowing against your own ability to think, remember, and create, and the bill comes due later, usually in an exam room with no Wi-Fi."

ASK: "Has anyone here ever asked ChatGPT to write something for school, used it, and then realized afterward you couldn't actually explain what you 'wrote'?" *[This usually gets an honest laugh and a few honest hands. Don't shame anyone — normalize it, then pivot.]*

"No judgment — basically everyone in this room has done some version of that, including adults at work. The goal today isn't to make you feel guilty. It's to make you *aware*, so you can choose differently."

2.4 Easy, relatable examples of AI putting your skills at risk

"Quick list — tell me if any of these sound familiar:

- **Letting AI finish your sentence in every caption or comment** — so you stop developing your own voice and writing style.
- **Asking AI to summarize a book or article you were assigned, instead of reading it** — you walk into class with information, but zero understanding.
- **Using AI to solve every math problem step-by-step instead of struggling with it first** — you get the right answer on the homework, and the wrong answer on the exam, because the struggle is where the learning actually happens.
- **Believing an AI-generated image or video instantly, because it looks real.**

On that last one — this is now a real, growing danger, not just a school problem. In one widely reported case, a woman in France was scammed out of **850,000 dollars** over 18 months by someone using AI-generated images and videos to convincingly pretend to be the actor Brad Pitt. Ads have circulated on Meta platforms using AI-cloned celebrity voices over real footage to push fake giveaways — Taylor Swift's voice was cloned to promote cookware she never endorsed. People have lost over a million dollars to deepfake videos of Elon Musk promoting fake crypto investments. These aren't edge cases anymore — fraud researchers estimate deepfake content online jumped from around half a million pieces in 2023 to roughly 8 million by 2025.

The common thread in every one of those scams? The victim's brain saw a familiar, trusted face and **skipped the step of questioning it**. That's critical thinking switching off at the exact moment it mattered most."

PART 3 — Who's Really Making the Decision? AI, Ads, and Your Choices (9–10 min)

"Let's go back to that Instagram feed from the beginning. Why did it show you exactly what it showed you?"

3.1 Recommender systems, in plain language

"Every time you scroll, like, skip fast, or rewatch something, you're feeding a system data. That system — called a **recommender system** — has one job: keep you watching, and slot in things you're likely to buy or click. It's not magic. It's pattern-matching, the same prediction trick from Part 1, just applied to *your behavior* instead of words."

3.2 ChatGPT now sells ads too

[Slide: ChatGPT ad screenshot — sponsored box under an answer]

"Here's something brand new, and most adults don't even know about it yet. As of February 2026, OpenAI started showing **ads inside ChatGPT itself** for free users in the US. If you ask ChatGPT about, say, running shoes, you might get your normal answer — and then a small labeled 'Sponsored' box underneath, suggesting a specific shoe brand that paid to be there. OpenAI says the ads won't change the actual answer's content, and they currently won't show ads to anyone under 18 — but the principle still matters: **the company that built the 'smart assistant' you trust for honest answers is also the company being paid by advertisers to put suggestions in front of you.** That's not necessarily evil — it's literally how Google Search and Instagram have made money for years — but it means every time you ask an AI 'what should I buy,' you should ask yourself: *who benefited from the answer I just got?*"

3.3 AI is also listening to what you say *about* brands

"It goes the other direction too. Companies now use AI to read through thousands — sometimes millions — of social media comments, reviews, and tweets to figure out how people genuinely feel about them. Think about an app almost everyone in this room has used, like **Talabat**. Researchers have built AI systems specifically trained to read Egyptian Arabic dialect — not just formal Arabic — to scan customer reviews and detect complaints, praise, and frustration at scale, far faster than any human team could. The company then uses that to fix problems, retrain staff, or completely redesign an ad campaign to hit the exact emotional nerve that's trending in the comments.

So while you're venting in a comment thread about bad delivery or commenting that you love a product — that text might directly become training data for the next ad campaign aimed at people just like you."

3.4 Egypt is part of this, at scale

"This isn't some far-off Silicon Valley thing. Egypt's digital ad market is already worth close to **1.8 billion US dollars** a year, and growing. Egypt has its own official **National AI Strategy** running through 2030, specifically because the government knows this technology is reshaping the economy. Local agencies in Cairo are already using AI tools to generate dozens of ad variations in hours instead of weeks, and to predict which ad is losing money before a human marketer even notices."

ASK: "Next time an ad feels *weirdly* relevant to you — like it's reading your mind — what's actually happening behind it?" *[Guide them to: it's not reading minds, it's pattern-matching your past behavior at scale, and it was probably tested against thousands of other people like you first.]*

PART 4 — AI in Music, Film, and Egyptian Marketing (5 min)

[Keep this section snappy and visual — this is the "fun" part before the serious closing message.]

"Quick tour of where else AI shows up in the stuff you consume for fun:

- **Music:** Tools like **Suno** can generate a full song — lyrics, vocals, instruments — from a single text prompt, in under a minute, in basically any language, including Arabic. People are already making full Mahraganat-style or pop tracks this way, no instrument or singing skill required.
- **Deepfakes in entertainment:** The same tech behind those scam videos also powers legitimate film and TV effects — de-aging actors, dubbing voices into other languages automatically, even resurrecting performances. The line between 'cool effect' and 'dangerous fake' is razor thin, and it's the *same underlying tool*.
- **Local Egyptian campaigns:** Agencies here are already using AI video tools that cut production time by 50–70%, generating an ad in 24–72 hours instead of the usual 2–6 weeks, and AI 'avatars' that can greet a customer by name in a personalized video — imagine an ad that's automatically slightly different for every single person who sees it."

ASK: "Can anyone name a song, ad, or video you've seen recently that you now suspect might have been AI-made — even partly?" *[Take a couple of guesses, validate the thinking process rather than the accuracy.]*

PART 5 — How to Actually Use AI (the part that matters most) (8–10 min)

"So after all that — scams, ad targeting, your brain quietly getting weaker — am I telling you to delete ChatGPT and never touch AI again?"

No. That would be like telling you to never use a calculator. AI isn't going away, and refusing to learn it isn't going to protect you — it's just going to leave you behind people who *do* learn to use it well. The goal is one simple rule:

AI should be the executor, never the decision-maker. You decide. AI does the boring, repetitive part. You verify everything it gives you.

Let's make that concrete."

5.1 Three relatable examples of using AI *correctly*

1. **Studying for an exam:** Bad use — "AI, answer these 50 past-paper questions for me." Good use — "AI, here are my own messy notes from class. Organize them into a clean summary and quiz me with 10 questions." You did the learning; AI did the boring formatting and repetition.
2. **A school project:** Bad use — "AI, write my whole essay on the Egyptian revolution." Good use — "I already know my argument and I've picked my three points. AI, help me check my grammar and suggest where my paragraph order is confusing." You made the decisions; AI executed cleanup work and you read every word it changed before submitting.
3. **Dealing with something that looks suspicious online** — a "celebrity" DM-ing you about a giveaway, a video that seems off. Bad use — believing it because it *looks* real. Good use — pause, and actually use AI as a tool to help you *verify*: ask it what the known patterns of this exact scam type look like, then go check the celebrity's verified account yourself before deciding anything. AI helps you think faster — it doesn't think *for* you.

ASK: "Give me one task from your own life — school or otherwise — where AI doing the boring part, while *you* stay in charge of the decision, would actually help you?" *[Take 2–3 genuine answers. This is the moment they personalize the whole lecture.]*

5.2 The closing reveal — proof, not just a promise

"I want to end with something honest. Preparing today's lecture, I needed to gather real, up-to-date examples — the deepfake cases, the exact ChatGPT ads details, the MIT study, Egypt's ad market numbers — accurately and fast. I used an AI model, Claude, to help research and draft a first version of this script.

[Optional: show this on a slide as proof — this is the actual instruction I gave it.]

"" Hi, I am preparing for a lecture about AI for students aged 15 to 18. I want a script I can follow that is engaging and interacts with the students.

The title of the lecture is "Who's Playing Who?" (من يلعب على مين؟), and it covers four main areas: Technology — what AI is and whether AI actually thinks; how we lose creativity and critical thinking because of it; how AI influences our decisions through ads and recommendations; and AI in art — music, Instagram profiles, and deepfakes.

Point 1 — What is AI: AI takes many different shapes and forms — for example, social media and how AI builds our feed, plus other parts of daily life, with examples given. Then introduce the main topic: LLMs, and the big LLM companies — Google, Anthropic, and OpenAI — who created Gemini, Claude, and ChatGPT respectively. After that, explain how LLMs work using the example of WhatsApp's keyboard autocomplete, and how LLMs essentially memorize patterns. Then ask: does AI actually think? Use a screenshot of ChatGPT showing "Thinking..." before it answers, and explain that LLMs mimic the behavior of thinking but do not actually think.

Point 2 — Critical thinking and creativity: Explain that one of humanity's greatest strengths is critical thinking ability. First explain what critical thinking is, using simple example sentences and asking the students to spot what's wrong with them — keep the examples obvious and easy to explain. Then introduce the "white paper" problem to showcase the creativity aspect. Then give easy, understandable examples of bad uses of LLMs that put students' critical thinking and creativity skills at risk.

Point 3 — Ads and recommendations: Explain how people and companies use AI to deliver ads and recommendations, with a short section on recommender systems in daily life and how they can influence what users buy or use. Mention that OpenAI has an ads option, where companies that pay can get ChatGPT to give certain recommendations. Also mention how companies use AI to analyze people's conversations and opinions on social media about their companies or products, in order to generate better content, ads, or brand positioning to increase sales. Ground everything in simple, understandable examples.

Point 4 — AI in creative industries: Keep this part small — just mention the use of AI in music, film, and TV ad production, and how local brands in Egypt generate marketing campaigns using AI. Talk about Suno and any other relevant tools.

Point 5 — How to use AI properly: Explain how to use AI and LLMs properly — that AI should be the executor, not the decision-maker, and that everything it produces should be verified. LLMs were meant to do the dirty, repetitive, boring work for us. Add simple, relatable examples. Finally, present this very prompt as a correct example of how to use AI, and how I personally apply what I am teaching.

Write the script in English, but remember — this is very important — the students are Egyptian boys, so ground all references and examples to them with real, reliable sources from the internet, and make the content focused and relatable to Egypt. ""

Notice what's in there: the structure, the five points, the order, the requirement that everything be real and Egypt-relevant, even the instruction to be honest with you about using AI right now. Every one of those decisions is mine. The AI's job was to go find the screenshots, the studies, the statistics, and arrange them — the boring digging.

But notice what I did, and didn't do. I decided the topic, the five sections, the message, and the tone for this exact room and your exact age group. I decided which examples were appropriate to bring to you and which weren't. I read every single fact it gave me and checked it before standing here and saying it to you. I am the one delivering it, adapting it live based on your reactions, and taking responsibility for every word.

The AI did the repetitive, boring digging — searching dozens of sources, organizing them, drafting a structure. I did the deciding, the verifying, and the actual thinking.

That's the whole lecture in one sentence: AI executes, you decide.

So — who's playing who? After today, that should be a question you can always answer for yourself."

Q&A / Closing discussion (remaining time)

Suggested closer questions if the room goes quiet:

- "What's one AI tool you use that you'll now use differently after today?"
- "Do you think AI companies *should* be allowed to put ads inside chatbot answers? Why or why not?"
- "If you could design one rule that every AI company had to follow, what would it be?"

Appendix — sources to double-check / cite on your slides

- MIT Media Lab, *"Your Brain on ChatGPT: Accumulation of Cognitive Debt when Using an AI Assistant for Essay Writing Task,"* Kosmyna et al., 2025 (media.mit.edu / arXiv 2506.08872)
- Stevenson, Smal, Baas, Grasman & van der Maas, *"Putting GPT-3's Creativity to the (Alternative Uses) Test,"* University of Amsterdam, 2022 (arXiv 2206.08932)
- OpenAI, *"Our approach to advertising and expanding access to ChatGPT,"* openai.com, Jan 2026; CNN Business coverage of the ChatGPT ads launch, Jan 16 2026
- McAfee, *2025 Most Dangerous Celebrity: Deepfake Deception List*; Blackbird.AI deepfake celebrity scandal roundup; reporting on the Brad Pitt impersonation scam (~\$850,000, 18-month scam) and the Taylor Swift/Le Creuset Meta ad scam
- Egypt's National AI Strategy 2025–2030; Egypt digital ad spend figures (~\$1.79B projected 2025) — various Egyptian digital marketing industry reports, 2025–2026
- Suno AI (suno.com) — generative music tool capabilities
- Research on Egyptian-dialect Arabic sentiment analysis for food-delivery review platforms (PeerJ Computer Science, 2025)

Tip: swap in your own screenshot of ChatGPT's "Thinking..." animation, and pull a current screenshot of the ChatGPT "Sponsored" ad box before the lecture — both change often as these products update.